

Lab 5: Copy-on-write Fork

(100 points)

Due by the end of Tuesday, Nov. 25, 2025

Note: This is a group-based project.

Overview

Virtual memory provides a level of indirection: the kernel can intercept memory references by marking page table entries (PTEs) invalid or read-only, leading to page faults, and can change what addresses mean by modifying PTEs. There is a saying in computer systems that any system's problem can be solved with a level of indirection. This lab explores an example: copy-on-write fork.

Note: You should start with a clean codebase!

1. The Problem

The [fork\(\)](#) system call in xv6 copies all of the parent process's user-space memory into the child. If the parent is large, copying can take a long time. Worse, the work is often largely wasted: [fork\(\)](#) is commonly followed by [exec\(\)](#) in the child, which discards the copied memory, usually without using most of it. On the other hand, if both the parent and child use a copied page, and one or both write to it, the copy is truly needed.

2. The Solution

Your goal in implementing [copy-on-write \(COW\) fork\(\)](#) is to defer allocating and copying physical memory pages until the copies are actually needed, if ever.

[COW fork\(\)](#) creates just a page table for the child, with PTEs for user memory pointing to the parent's physical pages. [COW fork\(\)](#) marks all the user PTEs in both parent and child as **read-only**. When either process tries to write to one of these COW pages, the CPU will force a page fault. The kernel page-fault handler detects this case, allocates a page of physical memory for the faulting process, copies the original page into the new page, and modifies the relevant PTE in the faulting process to refer to the new page, this time with the PTE marked **writable**. When the page fault handler returns, the user process will be able to write its copy of the page.

[COW fork\(\)](#) makes freeing of the physical pages that implement user memory a little trickier. A given physical page may be referred to by multiple processes' page tables, and should be freed only when the last reference disappears. In a simple kernel like xv6 this bookkeeping is reasonably straightforward, but in production kernels this can be difficult to get right; see, for example, [Patching until the COWs come home](#).

3. Implement Copy-on-write Fork (80 pts)

Your task is to implement copy-on-write fork in the xv6 kernel. You are done if your modified kernel executes both the `cowtest` and '`usertests -q`' programs successfully.

To help you test your implementation, an xv6 user program called `cowtest.c` is provided. `cowtest` runs various tests, but even the first will fail on unmodified xv6. Thus, initially, you will see:

```
$ cowtest
simple: fork() failed
$
```

The "simple" test allocates more than half of the available physical memory, and then `fork()`s. The fork fails because there is not enough free physical memory to give the child a complete copy of the parent's memory.

When you are done, your kernel should pass all the tests in both `cowtest` and `usertests -q`. That is:

```
$ cowtest
simple: ok
simple: ok
three: ok
three: ok
three: ok
file: ok
forkfork: ok
ALL COW TESTS PASSED
$ usertests -q
...
ALL TESTS PASSED
$
```

Here's a reasonable plan of attack.

- Modify `uvmcopy()` in `kernel/vm.c` as `uvmcopy_cow()` to map the parent's physical pages into the child, instead of allocating new pages. Clear `PTE_W` in the PTEs of both child and parent for pages that have `PTE_W` set. Modify `kfork()` in `kernel/proc.c` to call `uvmcopy_cow()` instead of `uvmcopy()`.
- Modify `vmfault()` to recognize page faults. When a write page-fault occurs on a COW page that was originally writable, allocate a new page with `kalloc()`, copy the old page to the new page, and install the new page in the PTE with `PTE_W` set. Pages that were originally read-only (not mapped `PTE_W`, like pages in the text segment) should remain read-only and shared between parent and child; a process that tries to write such a page should be killed.
- Ensure that each physical page is freed when the last PTE reference to it goes away -- but not before. A good way to do this is to keep, for each physical page, a "reference count" of the number of user page tables that refer to that page. Set a page's reference count to one when `kalloc()` allocates it. Increment a page's reference count when fork causes a child to share the page, and decrement a page's count each time any process drops the page from its page table. `kfree()` should only place a page back on the free list if its reference count is zero. It's OK to keep these counts in a fixed-size array of integers. You'll have to work out a

scheme for how to index the array and how to choose its size. For example, you could index the array with the page's physical address divided by 4096, and give the array a number of elements equal to the highest physical address of any page placed on the free list by `kinit()` in `kalloc.c`. Feel free to modify `kalloc.c` (e.g., `kalloc()` and `kfree()`) to maintain the reference counts.

- Modify `copyout()` to use the same scheme as page faults when it encounters a COW page.

Some hints:

- It may be useful to have a way to record, for each PTE, whether it is a COW mapping. You can use the `RSW` (reserved for software) bits in the RISC-V PTE for this.
- `usertests -q` explores scenarios that cowtest does not test, so don't forget to check that all tests pass for both.
- Some helpful macros and definitions for page table flags are at the end of `kernel/riscv.h`.
- If a COW page fault occurs and there's no free memory, the process should be killed.

4. What to submit

In `lab5-firstname-lastname.txt`, please include the time logging for this lab (10 pts). Note that each group member should have their own time logging. In `README-lab5-lastname1-lastname2.txt`, write down a few words of explanation for each of the files you had to modify/create (10 pts). The three text files must be included in your source directory.

AFTER committing all changes, you run `make clean`, and `tar zcf .../lab5-lastname1-lastname2.tar.gz .../xv6lab5` assuming your xv6 source directory is `xv6lab5`. Then one group member uploads the complete tarred and compressed xv6 directory (`lab5-lastname1-lastname2.tar.gz`) on Canvas for the group. **Each group member needs to finish a peer evaluation form and upload it on Canvas.**